



## **Hadoop: Hands-On: Running MapReduce**

## Running a MapReduce Job:

**Command 1:** Create a directory in HDFS to store your input file

```
Shdfs dfs -mkdir /user/training/inputfolder
```

**Command 2:** Upload the file from the local directory to the input folder in HDFS

```
Shdfs dfs -put <your local input file address> /user/training/inputfolder
```

**Command 3:** Run your MapReduce job and specify an output folder in HDFS to store the output

```
Shadoop jar WordCount.jar WordCount /user/training/inputfolder  
/user/training/outputfolder
```

**Command 4:** Check the status of the job using the YARN UI in the web browser

Enter 'localhost:8088' in the URL bar if accessing through VM.  
Enter '<public IP of the server>:8088' in the URL bar if accessing remotely.

